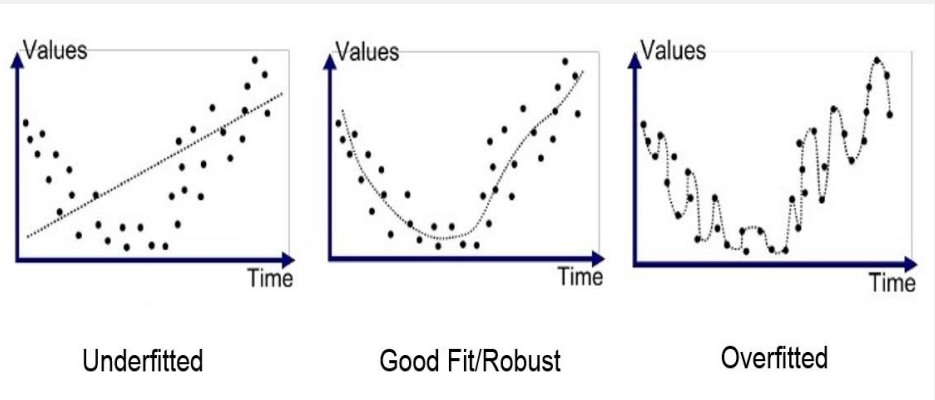




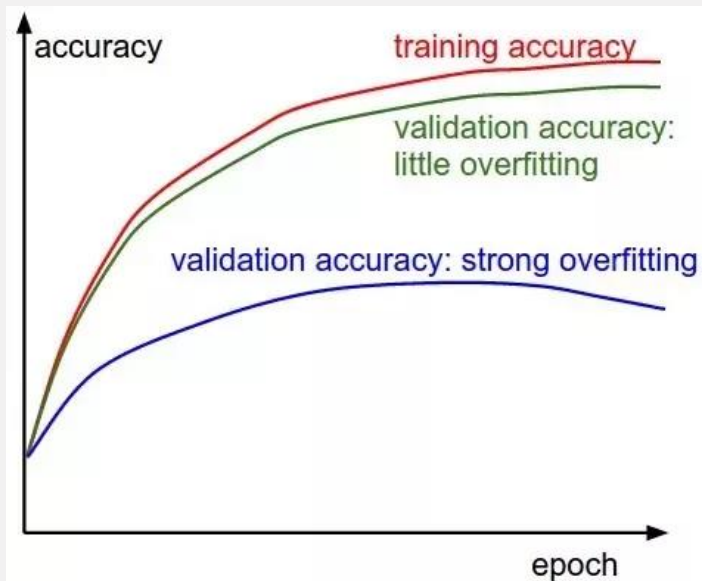
# TENSOR DECOMPOSITION METHOD COMPARISON

By: Leonel Ramirez, Alan Urteaga, Miguel Gutierrez, Het Patel



# Problems with CNNs

- CNNs are suited for learning these intricate & complex relationships. However, this comes with added risk of overfitting and lack robustness.
- Model generalizability is needed.
- DNNs leverage the spatial structure of input data via convolutions, pooling, point-wise non-linearities (like ReLU), etc. However, this structure typically wasted by a flattening layer followed by one or several fully connected layers.
- Tensor dropout provides a hand in this.



# What is a Tensor?

- **Tensor** - can be described as an  $n$ th-dimensional array
- **Mode** - is the dimension of the tensor or tensor shape.  
**Ex:** a tensor of mode 3 will be have a 3D shape and a mode 2 tenor will have a 2D shape
- **Rank** - is the number of underlining components that the tensor can be broken down into when computing tensor decomposition.

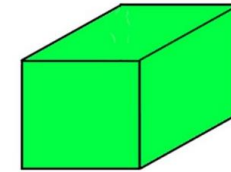
1D TENSOR /  
VECTOR

5
7
4 5
1 2
-6
3
2 2
1
6
3
-9

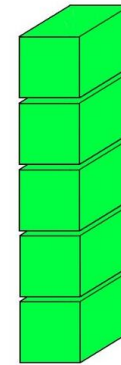
2D TENSOR /  
MATRIX

-9	4	2	5	7
3	0	1 2	8	6 1
1	2 3	-6	4 5	2
2 2	3	-1	7 2	6

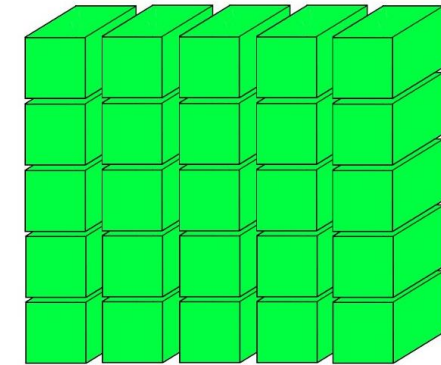
3D TENSOR /  
CUBE



-9	4	2	5	7
3	0	1 2	8	6 1
1	2 3	-6	4 5	2
2 2	3	-1	7 2	6



4D TENSOR  
VECTOR OF CUBES



5D TENSOR  
MATRIX OF CUBES

# Tensor Dropout

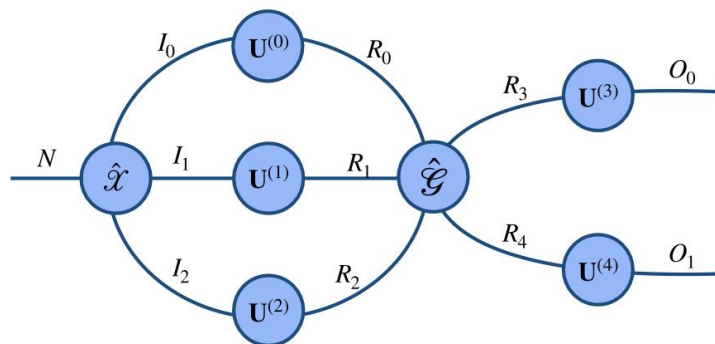
- **Stochastic**
- **Decomposed Forms**
- **Applied to Tucker and CP**

*Theorem 1:* Tensor Dropout on Tucker decomposition and deterministic regularized loss (proof in Appendix<sup>1</sup>).

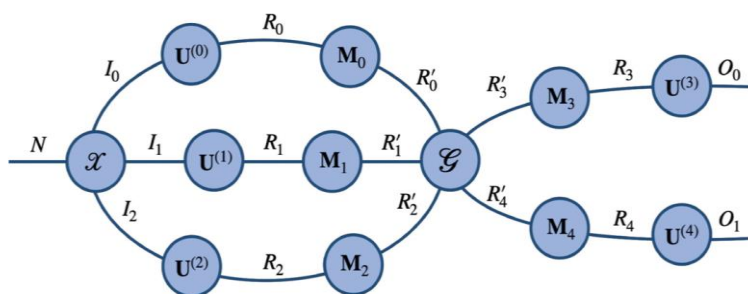
$$\begin{aligned}
 & \mathbb{E}_{\lambda} \left[ \frac{1}{S-1} \sum_{k=0}^{S-1} \left( \mathbf{y}^{(k)} - \langle \tilde{\mathcal{W}}, \mathcal{X}^{(k)} \rangle \right)^2 \right] \\
 &= \frac{1}{S-1} \sum_{k=0}^{S-1} \left( \mathbf{y}^{(k)} - \theta^N \langle \mathcal{W}, \mathcal{X}^{(k)} \rangle \right)^2 \\
 &+ \frac{\theta^N (1 - \theta^N)}{S-1} \sum_{k=0}^{S-1} \langle \mathcal{G}^{\star 2} \times_0 (\mathbf{U}^{(0)})^{\star 2} \dots \times_N (\mathbf{U}^{(N)})^{\star 2}, \\
 &(\mathcal{X}^{(k)})^{\star 2} \rangle. \tag{7}
 \end{aligned}$$

*Theorem 2:* Tensor Dropout on CP decomposition is equivalent to a deterministic regularized loss. (Proof in Appendix<sup>1</sup>)

$$\begin{aligned}
 & \mathbb{E}_{\lambda} \left[ \frac{1}{S-1} \sum_{k=0}^{S-1} \left( \mathbf{y}^{(k)} - \langle [\![\lambda; \mathbf{U}^{(0)}, \dots, \mathbf{U}^{(N)}]\!] , \mathcal{X}^{(k)} \rangle \right)^2 \right] \\
 &= \frac{1}{S-1} \sum_{k=0}^{S-1} \left( \mathbf{y}^{(k)} - \theta \langle [\![\mathbf{U}^{(0)}, \dots, \mathbf{U}^{(N)}]\!] , \mathcal{X}^{(k)} \rangle \right)^2 \tag{11} \\
 &+ \frac{\theta(1 - \theta)}{S-1} \sum_{k=0}^{S-1} \langle [\![(\mathbf{U}^{(0)})^{\star 2}, \dots, (\mathbf{U}^{(N)})^{\star 2}]\!] , (\mathcal{X}^{(k)})^{\star 2} \rangle.
 \end{aligned}$$



(a) Tensor diagram of a TRL



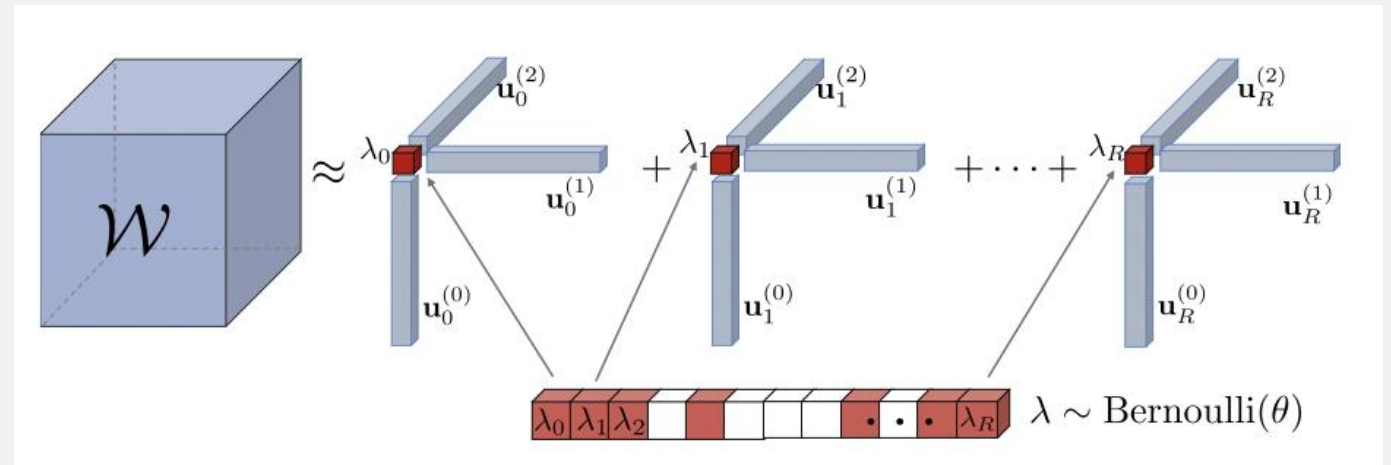
(b) Tensor diagram of a R-TRL

# Tensor Proposal

- Tensor Regression Networks propose to replace fully connected layers entirely with a tensor regression layer (TRL)
- This preserves the structure of the multidimensional data (input) by expressing an output tensor as the result of a tensor contraction between the input tensor and some low rank regression weight tensor
- TRL estimate the regression weight tensor  $\mathcal{W} = \mathbb{R}^{I_0 \times I_1 \times \dots \times I_N}$ , expressed under some low rank tensor decomposition

# Canonical-Polyadic (CP) Dropout Decomposition

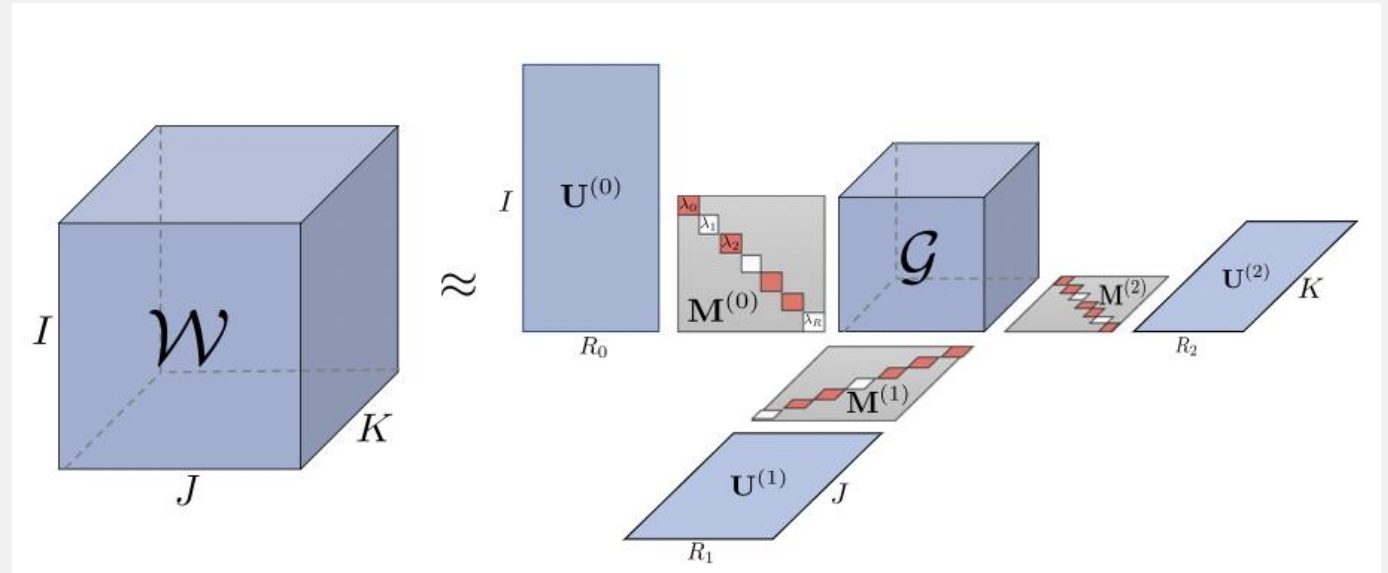
- Approximates the input tensor(weight) by summing the different vectors generated when decomposing the tensor (A, B, C)
- For the Drop out version of CP each component has a Bernoulli variable that determines if the component is dropped or not when decomposing the tensor. ("Acts like inducing stochasticity on the rank of decomposition.")[2]
- Vector sizes are of one of the mode of  $X$  by rank. Ex:  $X = R \times I \times J \times K$  is decomposed into A, B, C respectively. A.size =  $I \times r$



$$\mathcal{X} = \sum_{k=0}^{R-1} \underbrace{\lambda_k u_k^{(0)} \circ u_k^{(1)} \circ \dots \circ u_k^{(N)}}_{\text{rank-1 components}},$$

# Tucker Dropout Decomposition

- This method involves decomposing a tensor into n-modes, depending on its dimension.
- The result is a breakdown of a tensor into a low rank core  $\mathcal{G}$  (the dimensions are determined by the tensors rank), which projects the core along each mode with diagonal matrices (which are subject to *Bernoulli* tensor dropout) and set of factor matrices multiplied along the specified mode.



$$\mathcal{X} = \mathcal{G} \times_0 \mathbf{U}^{(0)} \times_1 \mathbf{U}^{(1)} \times \dots \times_N \mathbf{U}^{(N)}$$

$$\mathbf{M}^{(k)} = \text{diag}(\boldsymbol{\lambda}^{(k)})$$

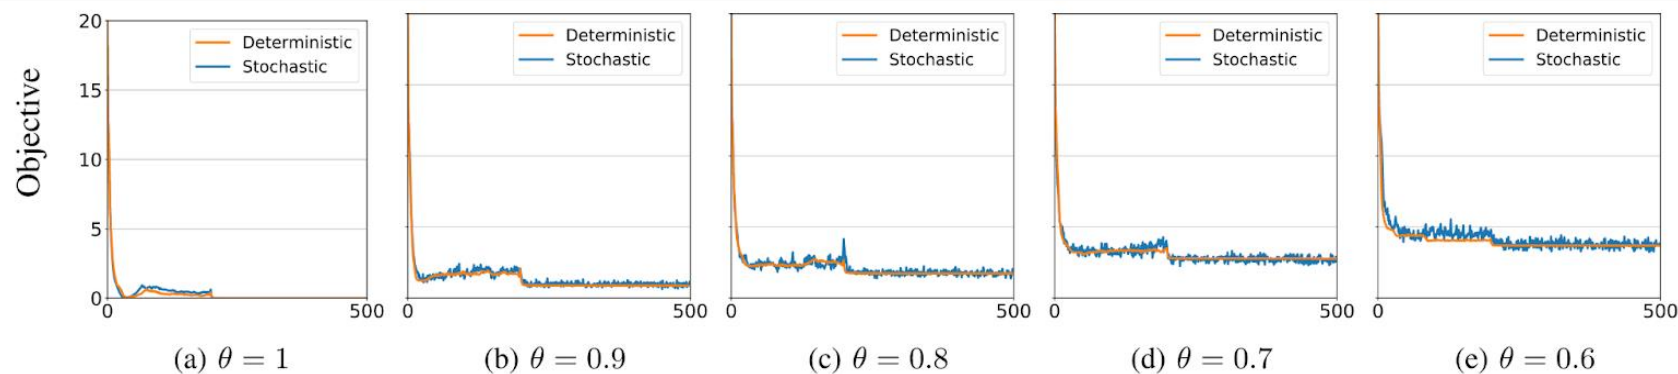


Fig. 4. **Experiment on synthetic data:** loss of the CP R-TRL as a function of the number of epochs for the stochastic version (orange) and the deterministic one based on the regularized objective function (blue). As expected, both formulations are empirically the same.

# 1. Synthetic Experiments

- 3rd Order Random Regression Weight Tensors
- Rank 10

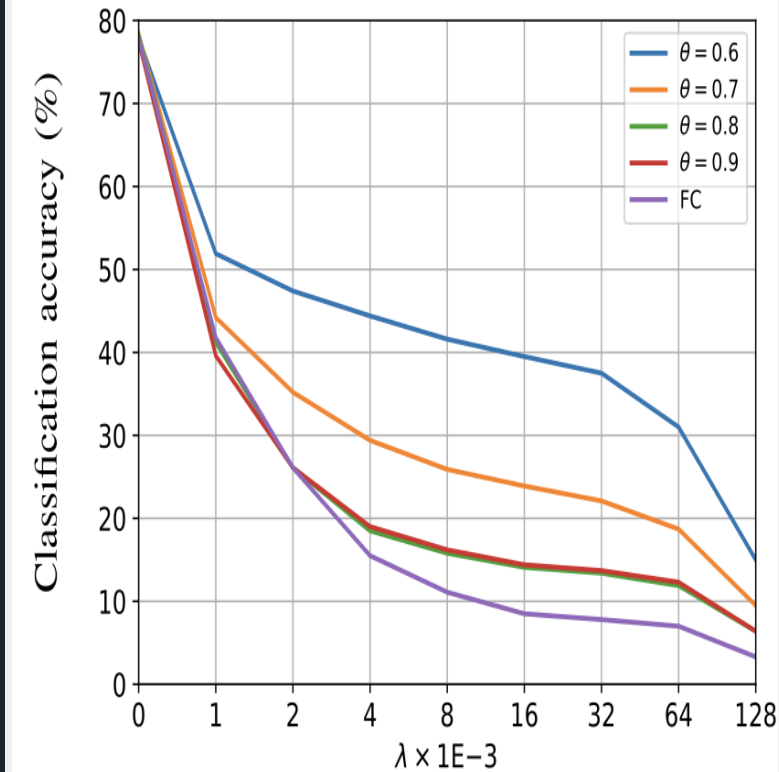


## 2. ImageNet

- **Robustness:** In addition to reducing overfitting, it is demonstrated that the proposed tensor dropout makes the model more robust to perturbations in input, for both random noise & adversarial attacks.
- The R-TRL architecture with tensor dropout is much more robust to adversarial attacks, depending on the tensor dropout rate ( $\theta$ ).

TABLE II  
REAL-VALUED NETWORK PERFORMANCE ON IMAGENET FOR FGSM, BIM  
AND PGD ATTACKS WITH  $\lambda \in \{2, 8, 16\}$ . WE REPORT CLASSIFICATION  
ACCURACY IN ALL CASES

Attack		Method			
Type	$\lambda$	Baseline	Ours		
			$\theta = 0.8$	$\theta = 0.7$	$\theta = 0.6$
Clean (no attack)		77.1	77.7	77.4	78.0
FGSM	2	26.1	26.0	35.3	47.4
	8	11.1	15.4	26.0	41.6
	16	8.5	14.1	24.0	39.5
BIM	2	26.1	26.0	35.3	47.3
	8	1.0	4.0	9.9	26.1
	16	0.1	1.0	2.8	13.2
PGD	2	26.0	26.0	35.6	47.1
	8	0.8	4.5	11.3	27.9
	16	0	1.4	3.8	13.5

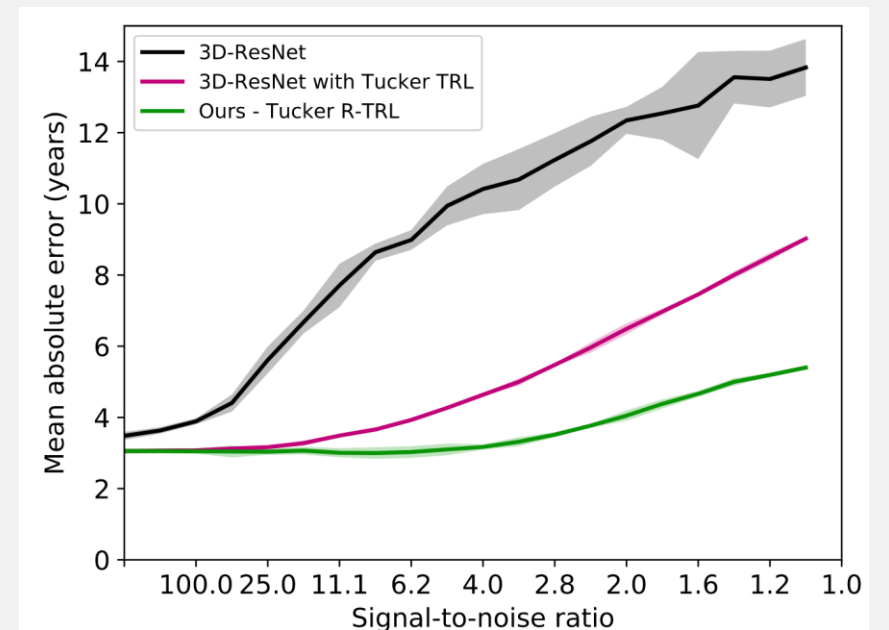


### 3. Phenotypic Trait MRI DATA

- Brain Age Difference
- Robustness Study: resistance to added gaussian white noise

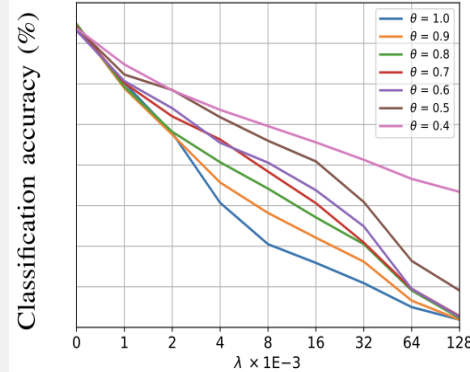
TABLE III  
CLASSIFICATION ACCURACY FOR U.K. BIOBANK MRI. THE RESNET MODELS WITH R-TRL SIGNIFICANTLY OUTPERFORMS THE VERSION WITH A FULLY-CONNECTED (FC) LAYER

Architecture	Regression	MAE
3D-ResNet	FC	2.96 years
3D-ResNet	Tucker	2.70 years
<b>Ours</b>	Randomized Tucker	<b>2.65 years</b>
<b>Ours</b>	Randomized CP	<b>2.58 years</b>

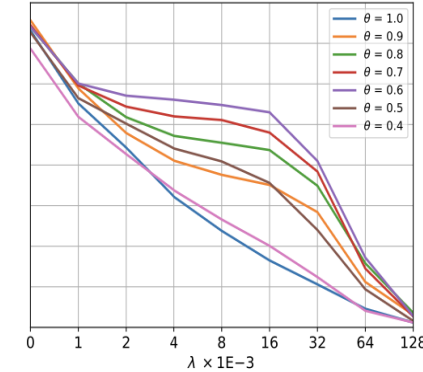


## 4. CIFAR-100

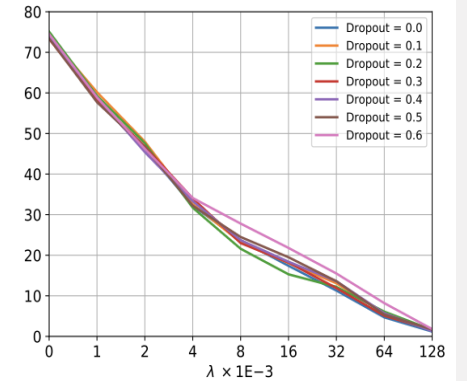
- Comparison of traditional tensor regression, and regular dropout, and asses the robustness of each method in the face of adversarial noise
- Tucker TRL + Tensor Dropout & CP TRL + Tensor Dropout prove to be more robust than Tucker TRL + Regular Dropout
- Results demonstrate that the performance is not overly sensitive to the choice of rank & tensor dropout probability



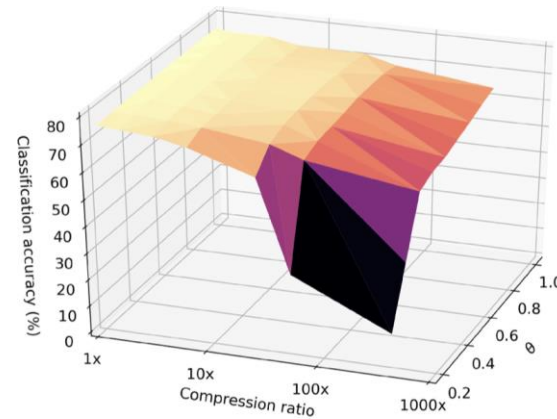
(a) Tucker TRL + Tensor Dropout. FGS attack on a Tucker TRL with Bernoulli tensor dropout and different drop rates.



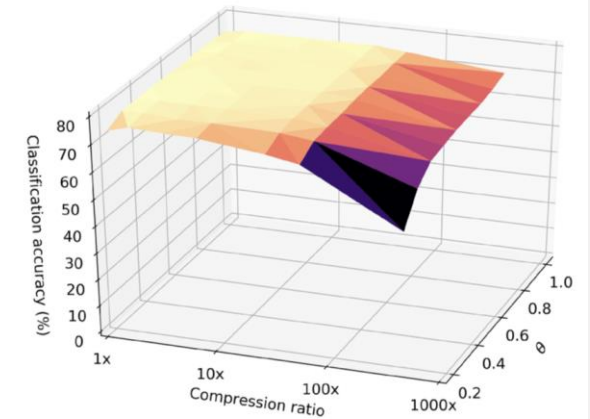
(b) CP TRL + Tensor Dropout. FGS attack on CP TRL with Bernoulli tensor dropout and different drop rates.



(c) Tucker TRL + Regular Dropout. FGS attack on Tucker TRL where *regular* dropout is applied to the regression weights, with different dropout rates.



(a) Tucker TRL + Tensor Dropout



(b) Tucker TRL + Tensor Dropout with replacement

# What is NORMO?

- Is an algorithm that approximates the number of components in a tensor decomposition
- Main purpose to reduce the number of components/rank (by determining the redundancy of certain components) to reduce computational cost
- Tested on artificial/real world tensor datasets with known number of components/rank using CP decomposition
- Performed the best on both high rank and low rank for both synthetic and real world datasets when compared to the different algorithms

**Table 4**

Estimation results in the real-world datasets with known model order (correct estimates in bold).

	True model Order	DIF FIT	CORC ONDIA	ARD ridge	ARD sparse	Convex Hull	N-D MDL	NOR MO
amino	4	3	<b>4</b>	◇	◇	[3 3 3]	11	<b>4</b>
dorrit	4	★	<b>4</b>	[13 6 13]	[10 5 12]	–	–	<b>4</b>
wbnmr	4	★	3	◇	◇	[24 25 25]	–	<b>4</b>
sugar	4	6	<b>4</b>	◇	◇	[1 4 4]	550	5
tongue	3	<b>3</b>	<b>3</b>	◇	◇	[2 2 1]	6	1

**Table 5**

Estimation results in the real-world datasets with unknown model order.

	DIF FIT	CORC ONDIA	ARD ridge	ARD sparse	Convex Hull	N-D MDL	NOR MO
enron	★	12	◇	◇	[2 2 2]	183	11
dblp	★	\	◇	◇	[25 25 7]	2722	†
challenge	★	7	◇	◇	\	123	5
friends	★	†	[25 25 17]	[25 25 12]	[25 25 16]	128	†
reality	★	11	◇	◇	[2,2,2]	92	11

Table from “A new method for estimating the number of components in CP tensor decomposition”, [1]

# Conclusions

- Paper [2] (Tensor Dropout) Tensor dropout helps constrain the learned representations to be more general, which leads to reduced overfitting & better out of sample generalization
- Spatial information is traditionally discarded during the flattening process, which is avoided using a tensor regression layer (TRL) and (R-TRL)
- Paper [1] (NORMO) proposed a different algorithm to determine the rank of a tensor decomposition along with accounting for the redundancy of some of the ranks
- Paper [3] (Tensor Methods) gives a detailed background on tensors and how they are used in computer vision and deep learning

# Next Steps weeks 3 - 5

- Fortify our knowledge weak points
- Become more familiar with the implementations of R-TRL from paper [2] (Tensor Dropout)
- Search for existing implementations for R-TRL from paper [2] (Tensor Dropout)
- Try to Replicate the results from paper [2] (Tensor Dropout)

**Thank you for  
your time!**

**Any Questions?**



# Citations:

- 1) Sofia Fernandes, Hadi Fanaee-T, João Gama, NORMO: A new method for estimating the number of components in CP tensor decomposition, Engineering Applications of Artificial Intelligence, Volume 96, 2020, 103926, ISSN 0952-1976, <https://www.sciencedirect.com/science/article/pii/S0952197620302530>
- 2) A. Kolbeinsson et al., "Tensor Dropout for Robust Learning," in IEEE Journal of Selected Topics in Signal Processing, vol. 15, no. 3, pp. 630-640, April 2021, doi: 10.1109/JSTSP.2021.3064182, <https://ieeexplore.ieee.org/document/9381098>
- 3) Y. Panagakis et al., "Tensor Methods in Computer Vision and Deep Learning," in Proceedings of the IEEE, vol. 109, no. 5, pp. 863-890, May 2021, doi: 10.1109/JPROC.2021.3074329, <https://ieeexplore.ieee.org/abstract/document/9420085>
- 4) "Tamara G. Kolda: 'Tensor Decomposition.'" YouTube. YouTube, April 6, 2018. <https://www.youtube.com/watch?v=L8uT6hgMt00>
- 5) "Applied Linear Algebra: Tensor Decompositions." YouTube. YouTube, December 2, 2020. <https://www.youtube.com/watch?v=tm5am60CIId4>
- 6) "Applied Linear Algebra: Implementing Tensor Decompositions." YouTube. YouTube, December 4, 2020. <https://www.youtube.com/watch?v=1UbAJeix3To>
- 7) dfleisch. "What's a Tensor?" YouTube. YouTube, November 20, 2011. <https://www.youtube.com/watch?v=f5liqUk0ZTw>.
- 8) Krishan. "Understanding Compression of Convolutional Neural Nets: Part 3." From Data to Decisions, June 9, 2020. <https://iksinc.online/2020/06/09/understanding-compression-of-convolutional-neural-nets-part-3/>.
- 9) Krishan. "Understanding Tensors and Tensor Decompositions: Part 1." From Data to Decisions, June 15, 2021. <https://iksinc.online/2018/02/12/understanding-tensors-and-tensor-decompositions-part-1/>
- 10) "Investigating Tensors with Pytorch." DataCamp. Accessed June 9, 2022. <https://www.datacamp.com/tutorial/investigating-tensors-pytorch>